

Ind Study W3

# Content

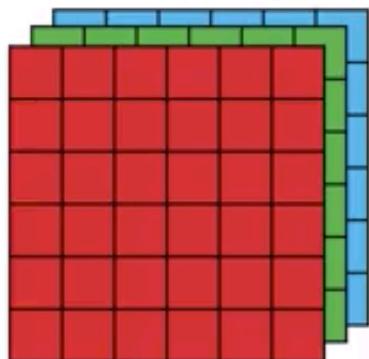
- CNN architecture
  - Alexnet
  - VGG
  - ResNet
  - GoogLeNet
- Transfer Learning
- Data Augmentation
- Multi-crop

# Alexnet

- Data Augmentation
  - Horizontal reflection
  - Altering the intensities of RGB channels in training images (Perform PCA)
- Dropout
  - Helps in coadoption of neurons in absence of some other neuron
- Contains local response normalization

## VGG - 16

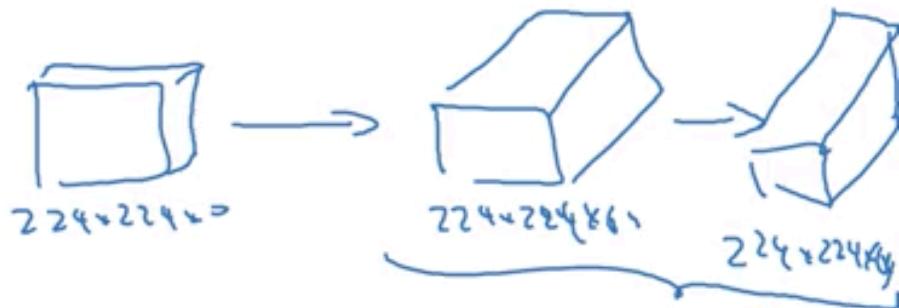
CONV =  $3 \times 3$  filter,  $s = 1$ , same



$224 \times 224 \times 3$

VAG-19

MAX-POOL =  $2 \times 2$ , s = 2



$\rightarrow 112 \times 112 \times 128 \rightarrow 56 \times 56 \times 128$

x2

[CONV 256]  
x3

$56 \times 56 \times 256 -$

POOL

$8 \times 256$

[CONV 512]  
× 3

POOL

$14 \times 14 \times 512$

[CONV 512]  
x3

$14 \times 14 \times 512$

POOL

×512 —

4096

→ FC  
4096

Softmax  
1000

$$n_H, n_w \downarrow$$

۸۲

~138m

# VGG

- Address depth of CNN( Fix all other parameters and increase depth)
- Preprocessing- Subtracting the mean RGB value
- Architecture
  - Filters- $3 \times 3$  and  $1 \times 1$ , stride-1
  - Same Padding -1 for  $3 \times 3$  filter
  - Followed by spatial + max pooling ( $f=2, s=2$ ) for some layers
  - Followed by 3 fully connected layers (same configuration in all network)
  - First two FC 4096 channels each , third 1000 channels
  - Final layer is softmax
  - Contains RELU but not local response normalization
- Having  $3 \times 3$  instead of  $7 \times 7$ :
  - Decreases number of parameters  $27C^2$  instead of  $49C^2$
  - Incorporate 3 non-linear rectification function instead of one, making decision function more discriminative
  - Imposing regularization

# VGG-Training

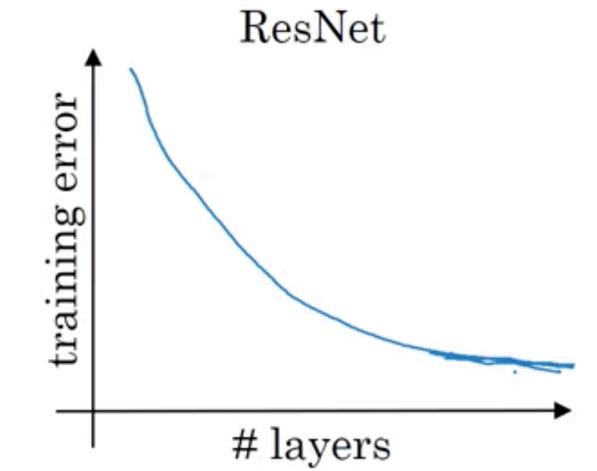
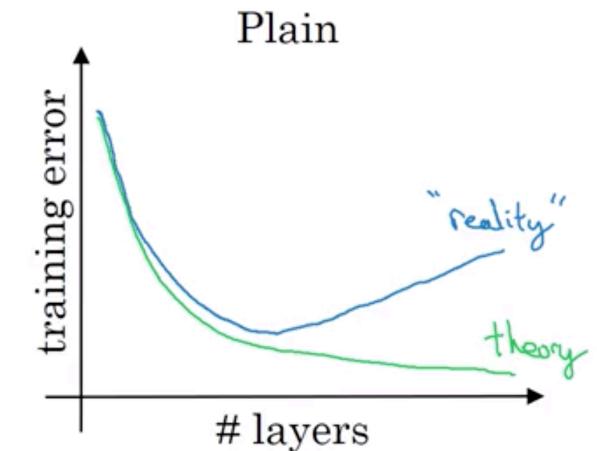
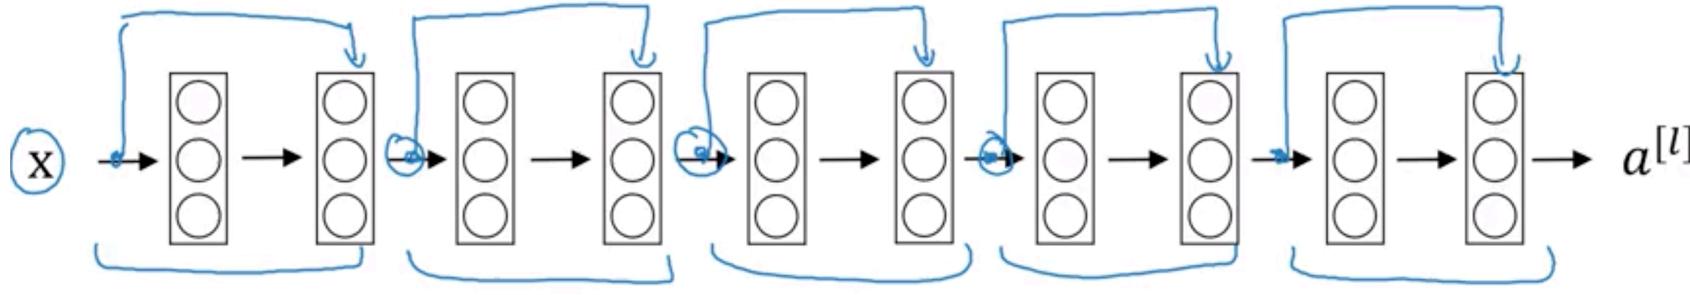
- Multinomial logistic regression with mini-batch gradient descent with momentum
- Batch size-256, momentum -0.9
- Weight decay and regularization (dropout ratio=0.5)
- Learning rate-  $10^{-2}$  , 370K iterations (74 epochs)
- Required less epochs due to
  - Implicit regularization imposed by greater depth and smaller conv. Filter sizes
  - Pre-initialization of certain layers
- Setting Training Image Size S: Rescaling training image by randomly sampling S from a certain range  $[S_{\min}, S_{\max}]$

# VGG- Experiments

- Single scale Evaluation
  - A deep net with small filters outperforms a shallow filters with larger filters
  - Scale jittering at training time ( $S \in [256,512]$  ) leads to significantly better results than training on images with fixed smallest side ( $S= 256$  or  $S=384$ ), even though a single scale is used at test time
- Multi-scale Evaluation
  - Running a model over several rescaled version of test image

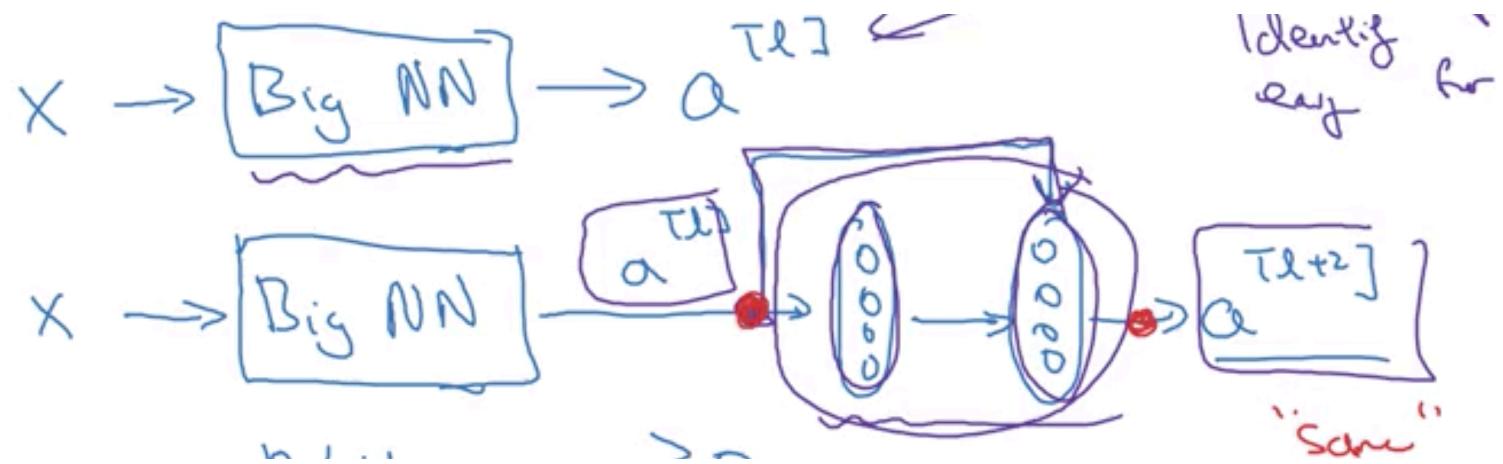
# ResNet

- Skipping layers to pass over information
- $a^{[l]} \rightarrow \text{Linear} \rightarrow \text{Relu} \rightarrow \text{Linear} \rightarrow \text{Relu} \rightarrow a^{[l+2]}$
- $a^{[l+2]} = g(z^{[l+2]} + a^{[l]})$ ;  $a^{[l]}$  passed to relu along with  $z^{[l+2]}$
- On plain network the training error tends to decrease first with increasing depth but increases further
- In ResNet error always decreases with depth



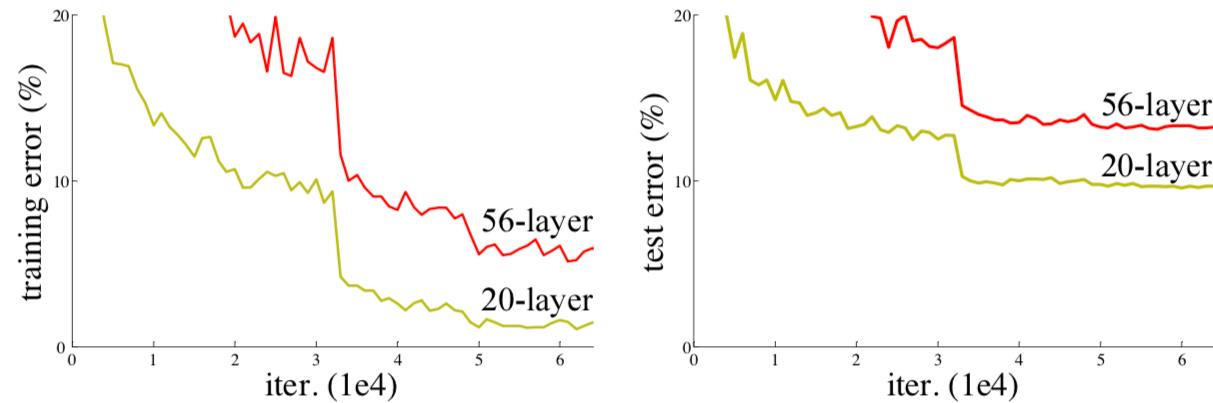
# ResNet

- $a^{[l+2]} = g(z^{[l+2]} + a^{[l]}) = a^{[l+2]} = g(w^{[l+2]} \cdot a^{[l+1]} + b^{[l+2]} + a^{[l]})$
- $a^{[l+2]} = g(a^{[l]})$ ; if  $w^{[l+2]} = 0, b^{[l+2]} = 0$
- Use same convolution to have dimension of  $z^{[l+2]}$  and  $a^{[l]}$  same



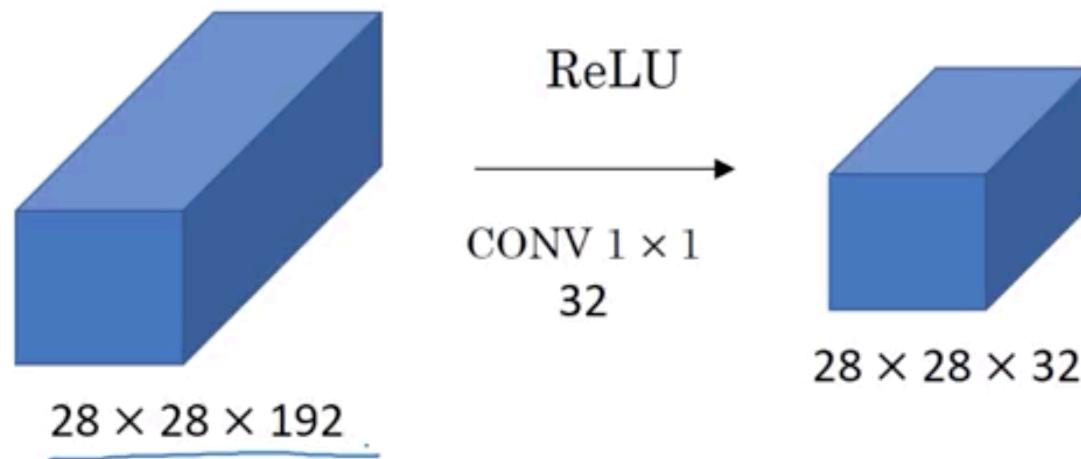
# ResNet

- Degradation with deeper network does not happen by overfitting
- Adding more layers to a deep model leads to higher training error
- Shortcut connections act as identity mappings
- ResNet converges faster than plain for 18 layers



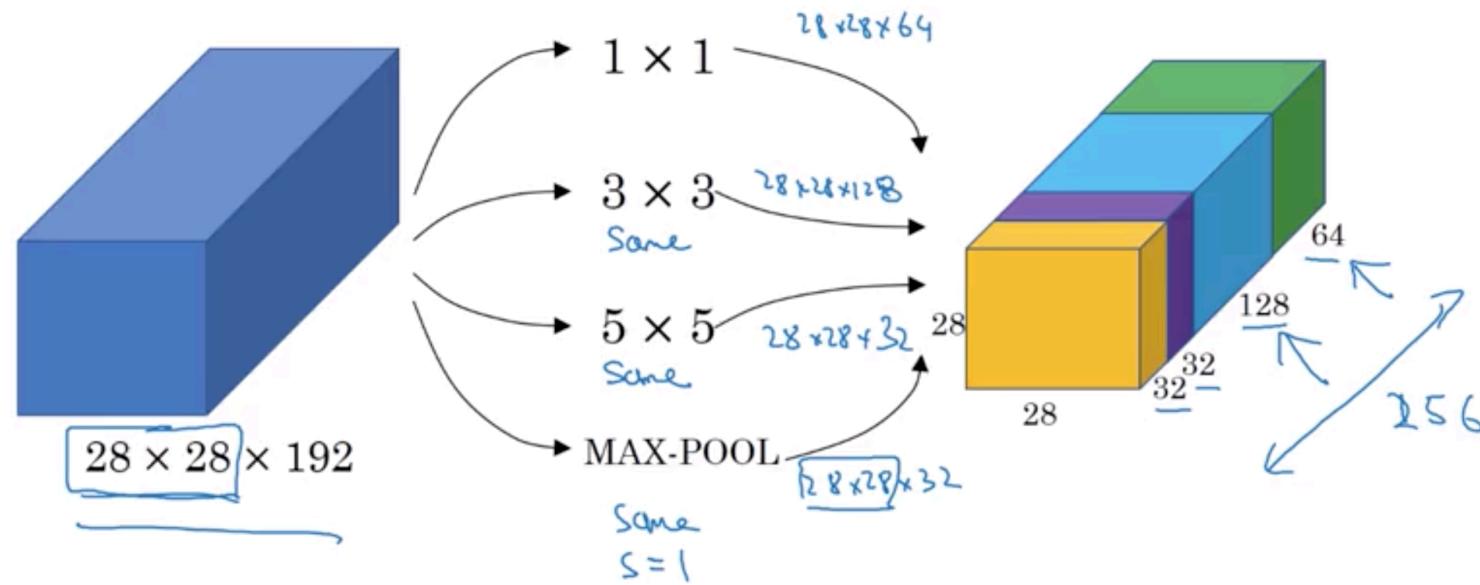
# Network in Network (1\*1 convolution)

- Multiply individual element with a number
- Helpful in shrinking number of channels and save computations



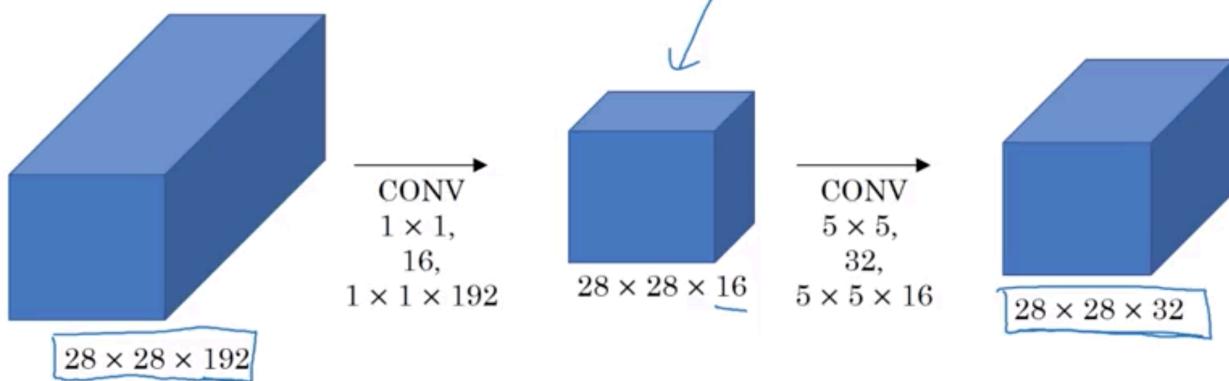
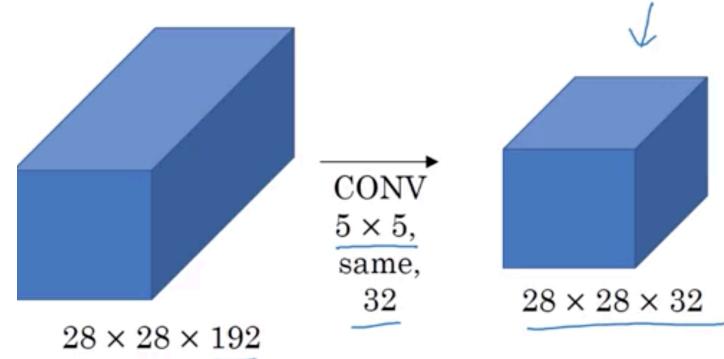
# GoogleNet( Inception)

- Try all type of convolutions  $1 \times 1$ ,  $3 \times 3$ ,  $5 \times 5$  and MAX-POOL
- Concatenate all the output and let network learn whatever parameters it want to use

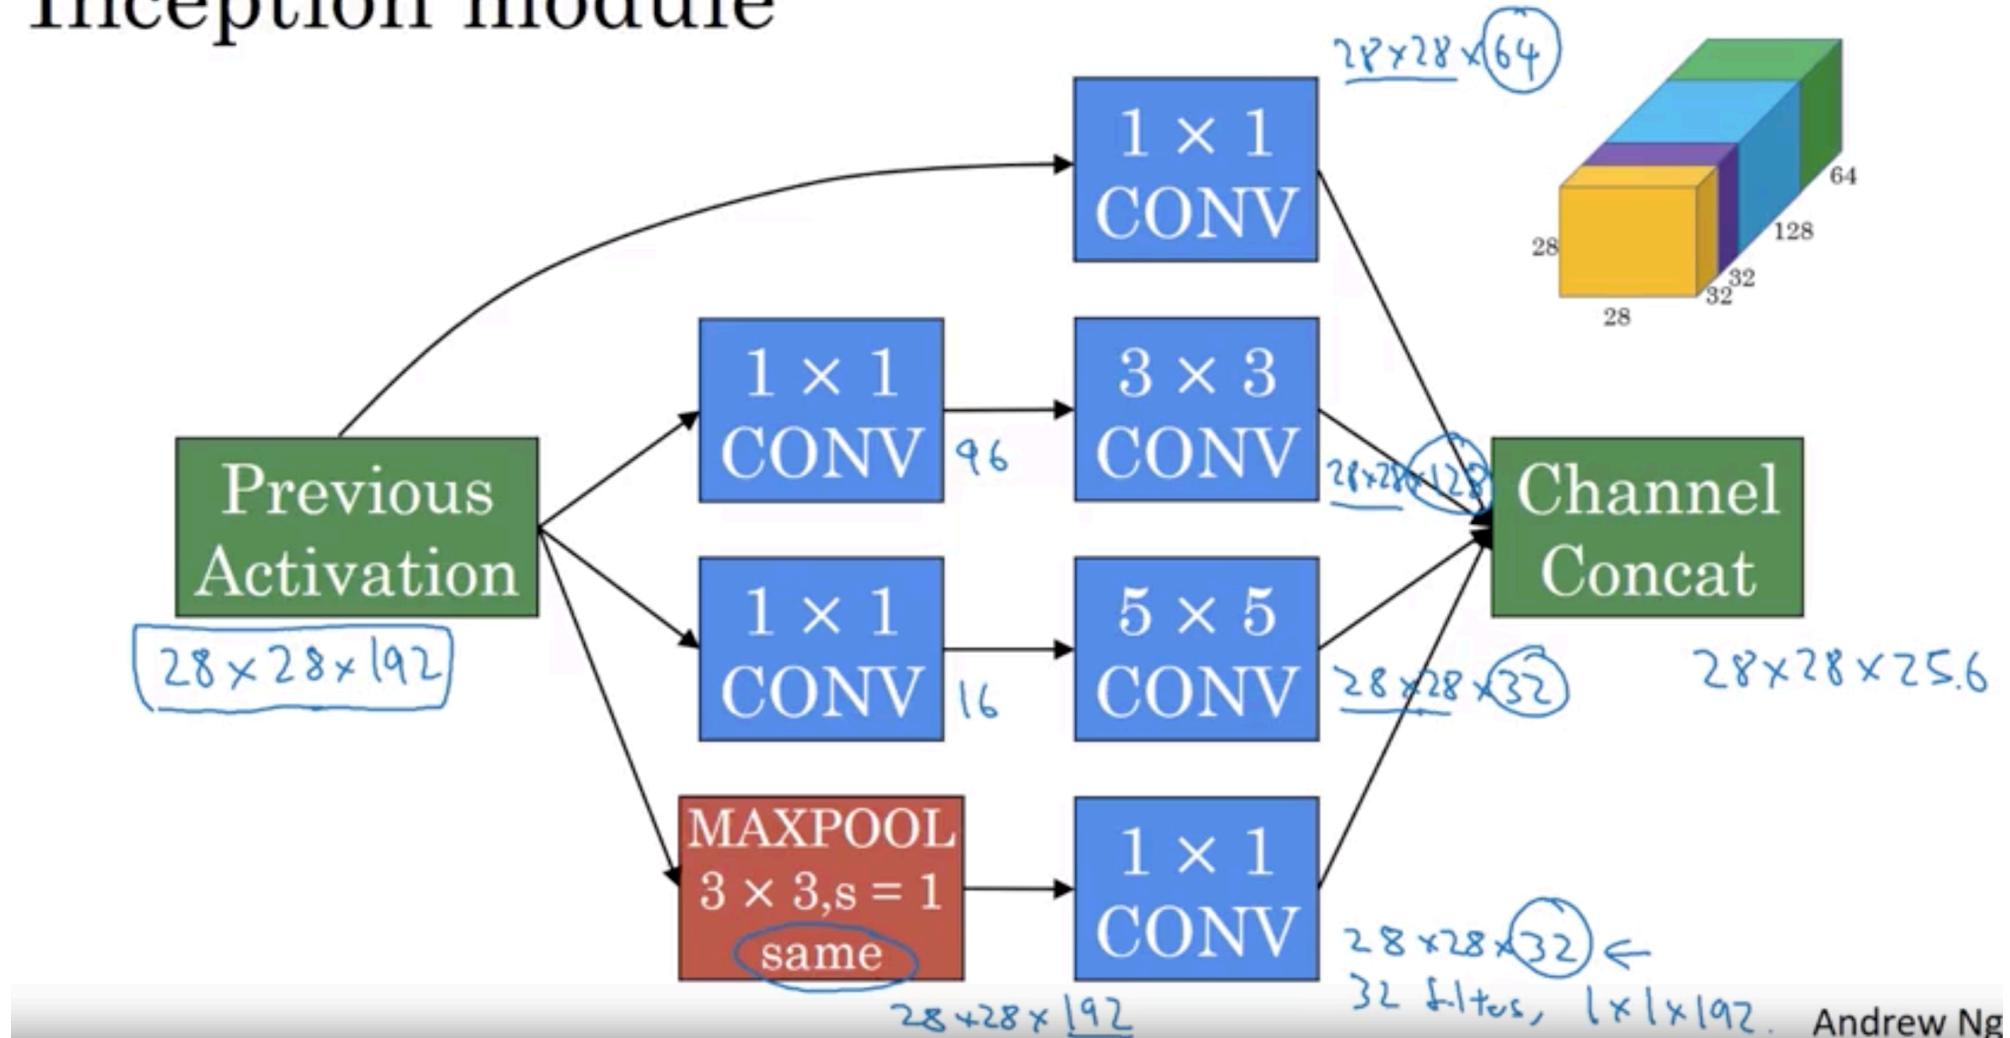


# Inception

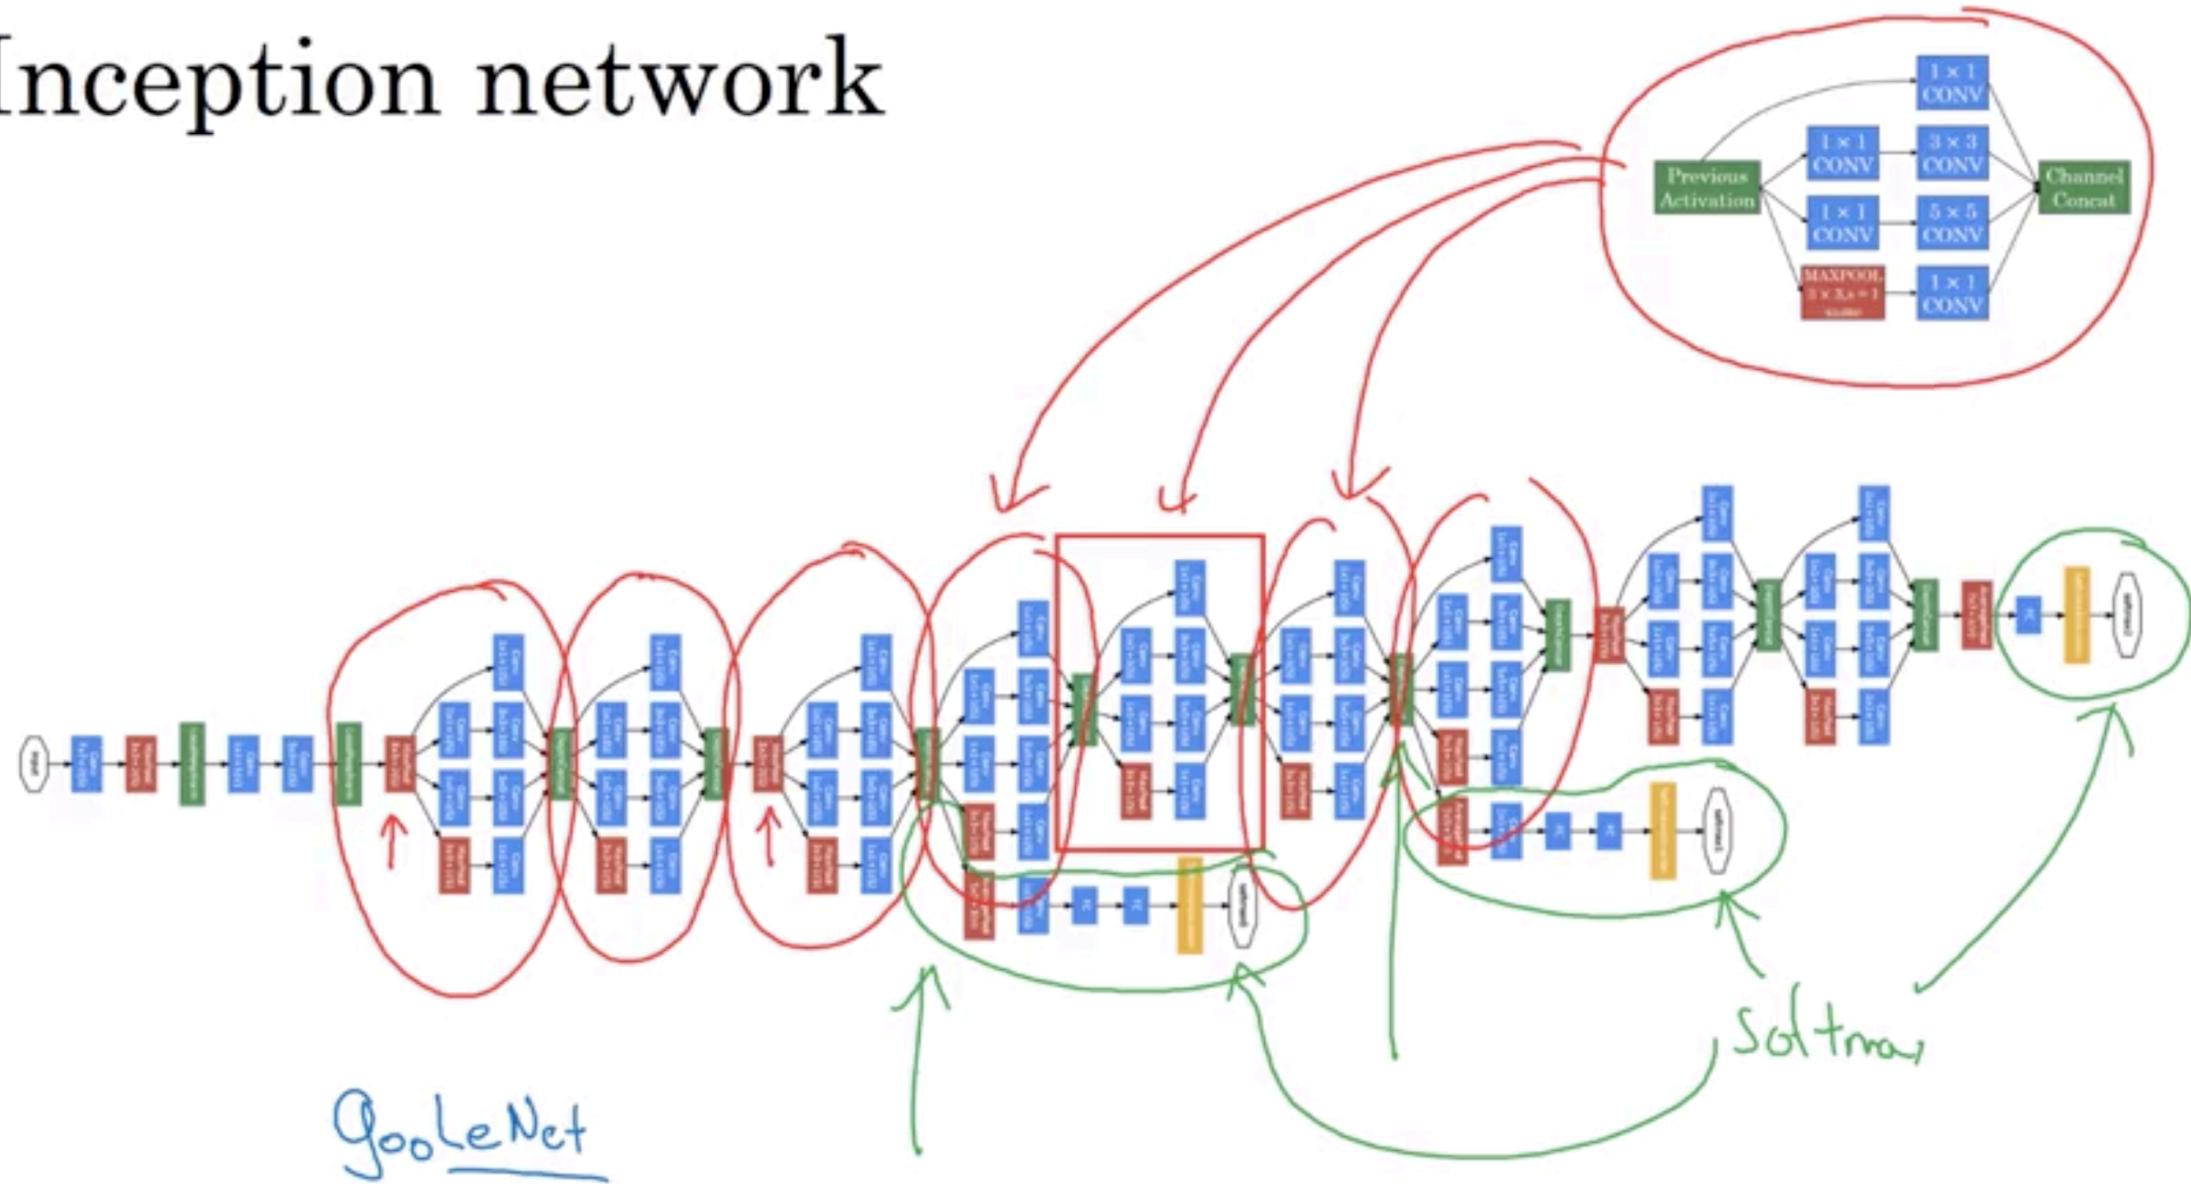
- Computation cost for 5\*5 CONV:
  - 32 filters of 5\*5\*192
  - $28*28*32* 5*5*192 = 120M$  (expensive)
  - Using 1\*1 convolution reduces cost by 1/10<sup>th</sup>
  - Introduce a bottleneck layer (28\*28\*16) in between
  - New cost= $28*28*16*1*1*192 + 28*28*32*5*5*16 = 12.4 M$



# Inception module



# Inception network



# Transfer Learning

- Reusing an existing model as starting point for other model
- It saves time and computation power which is required to build from scratch.

# Data Augmentation

- Mirroring
- Random Cropping
- Rotation
- Shearing
- Local Warping
- Color shifting- take different values of R, G, and B and use them to distort the color channels, bring more robustness with colors
- PCA color augmentation- if image is purple it mainly contains red and blue but lesser green, so it adds and subtracts more of RG and less of G

# Common Augmentation method

Mirroring

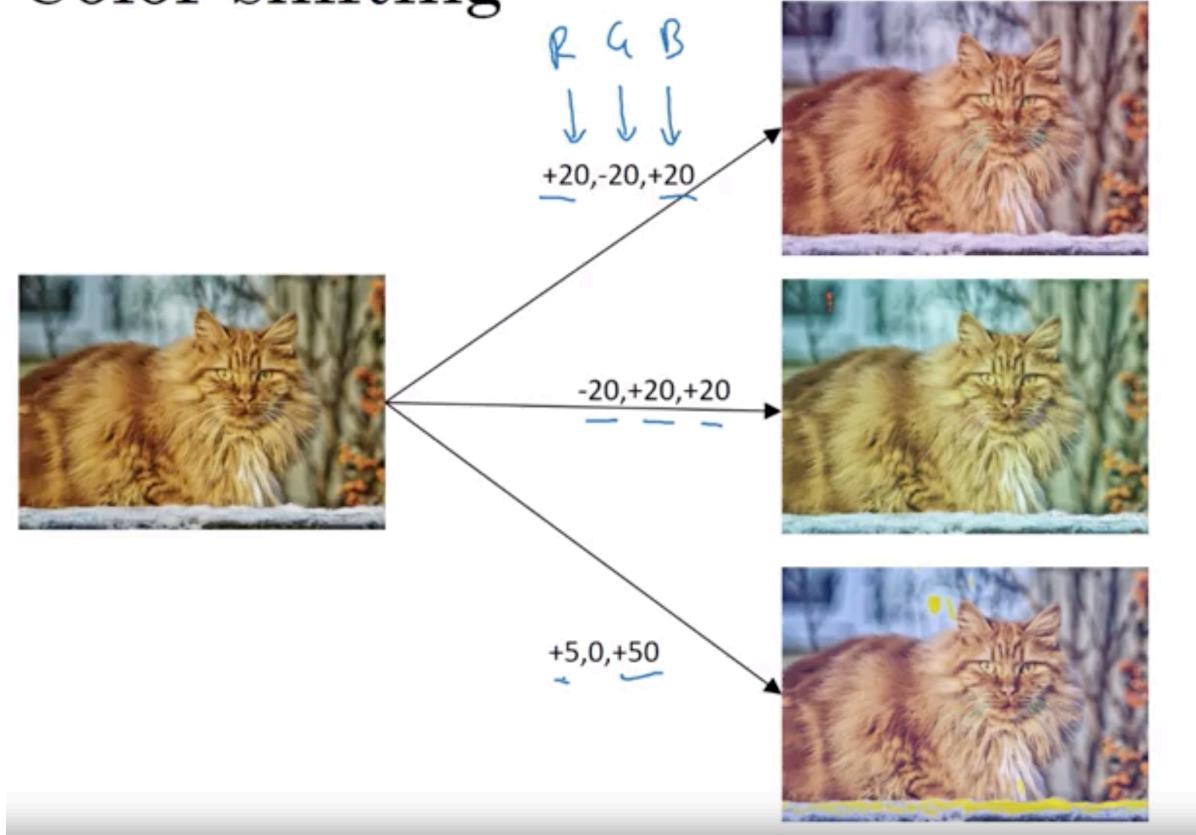


Random Cropping

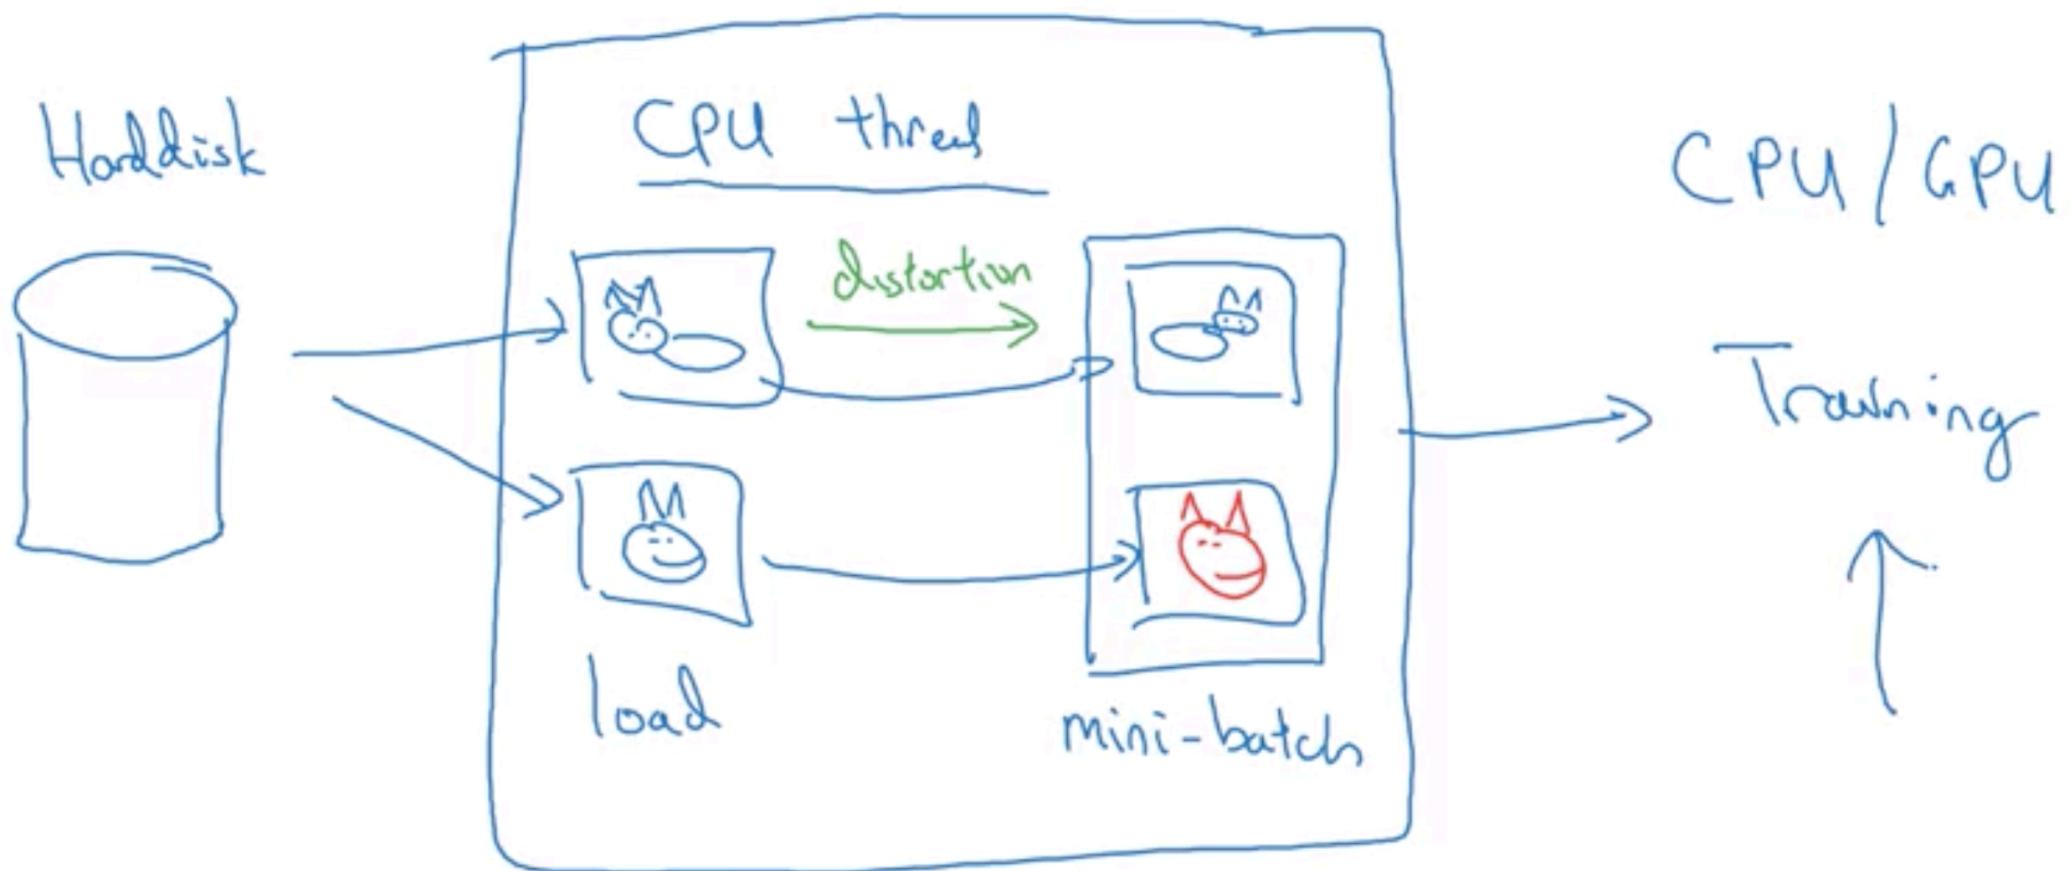


# Common Augmentation method

## Color shifting

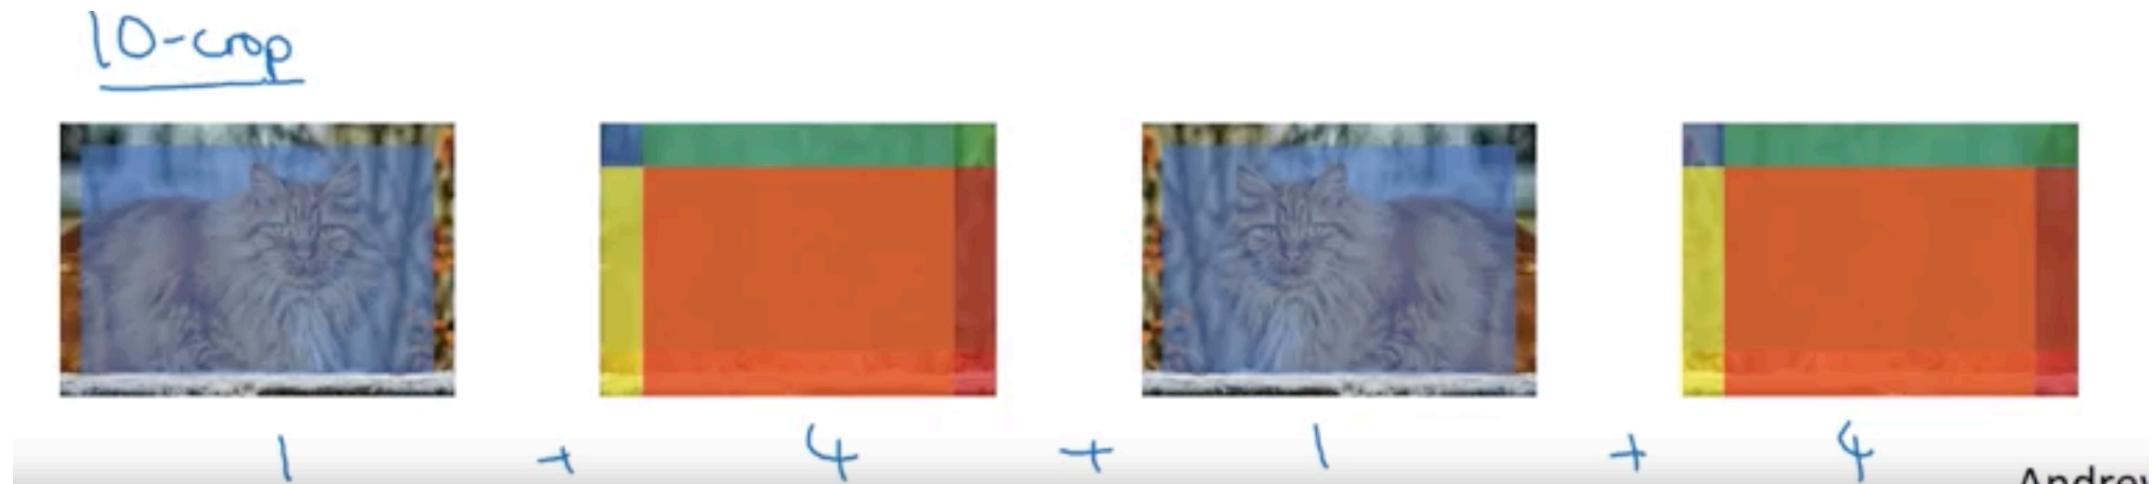


# Implementing distortions during training



# Multi crop at test time

- Run classifier on multiple versions of test images and average results
- 10 crop:



# Quiz

---



3. In order to be able to build very deep networks, we usually only use pooling layers to downsize the height/width of the activation volumes while convolutions are used with "valid" padding. Otherwise, we would downsize the input of the model too quickly.

0 / 1  
point



6. Which ones of the following statements on Residual Networks are true? (Check all that apply.)

0 / 1  
point



9. Which ones of the following statements on Inception Networks are true? (Check all that apply.)

0 / 1  
point



10. Which of the following are common reasons for using open-source implementations of ConvNets (both the model and/or weights)? Check all that apply.

0 / 1  
point

# References

- CUDA implementation of CNN:  
<https://code.google.com/archive/p/cuda-convnet/>
- Keras documentation: <https://keras.io/models/model/>
- ResNet Blog: <http://torch.ch/blog/2016/02/04/resnets.html>
- ResNet Github: <https://github.com/KaimingHe/deep-residual-networks>
- ResNet original paper: <https://arxiv.org/abs/1512.03385>
- Identity mapping in ResNet: <https://arxiv.org/abs/1603.05027>